

BPF+

Exploiting Global Data-flow Optimization in a Packet Filter Architecture

Andrew Begel, Steven McCanne, Susan L. Graham

University of California, Berkeley

with acknowledgments to Van Jacobson

The Big Picture

With BPF+, we can express packet filters
in a high-level language *and* compile
them into efficient code.

The Packet Filter

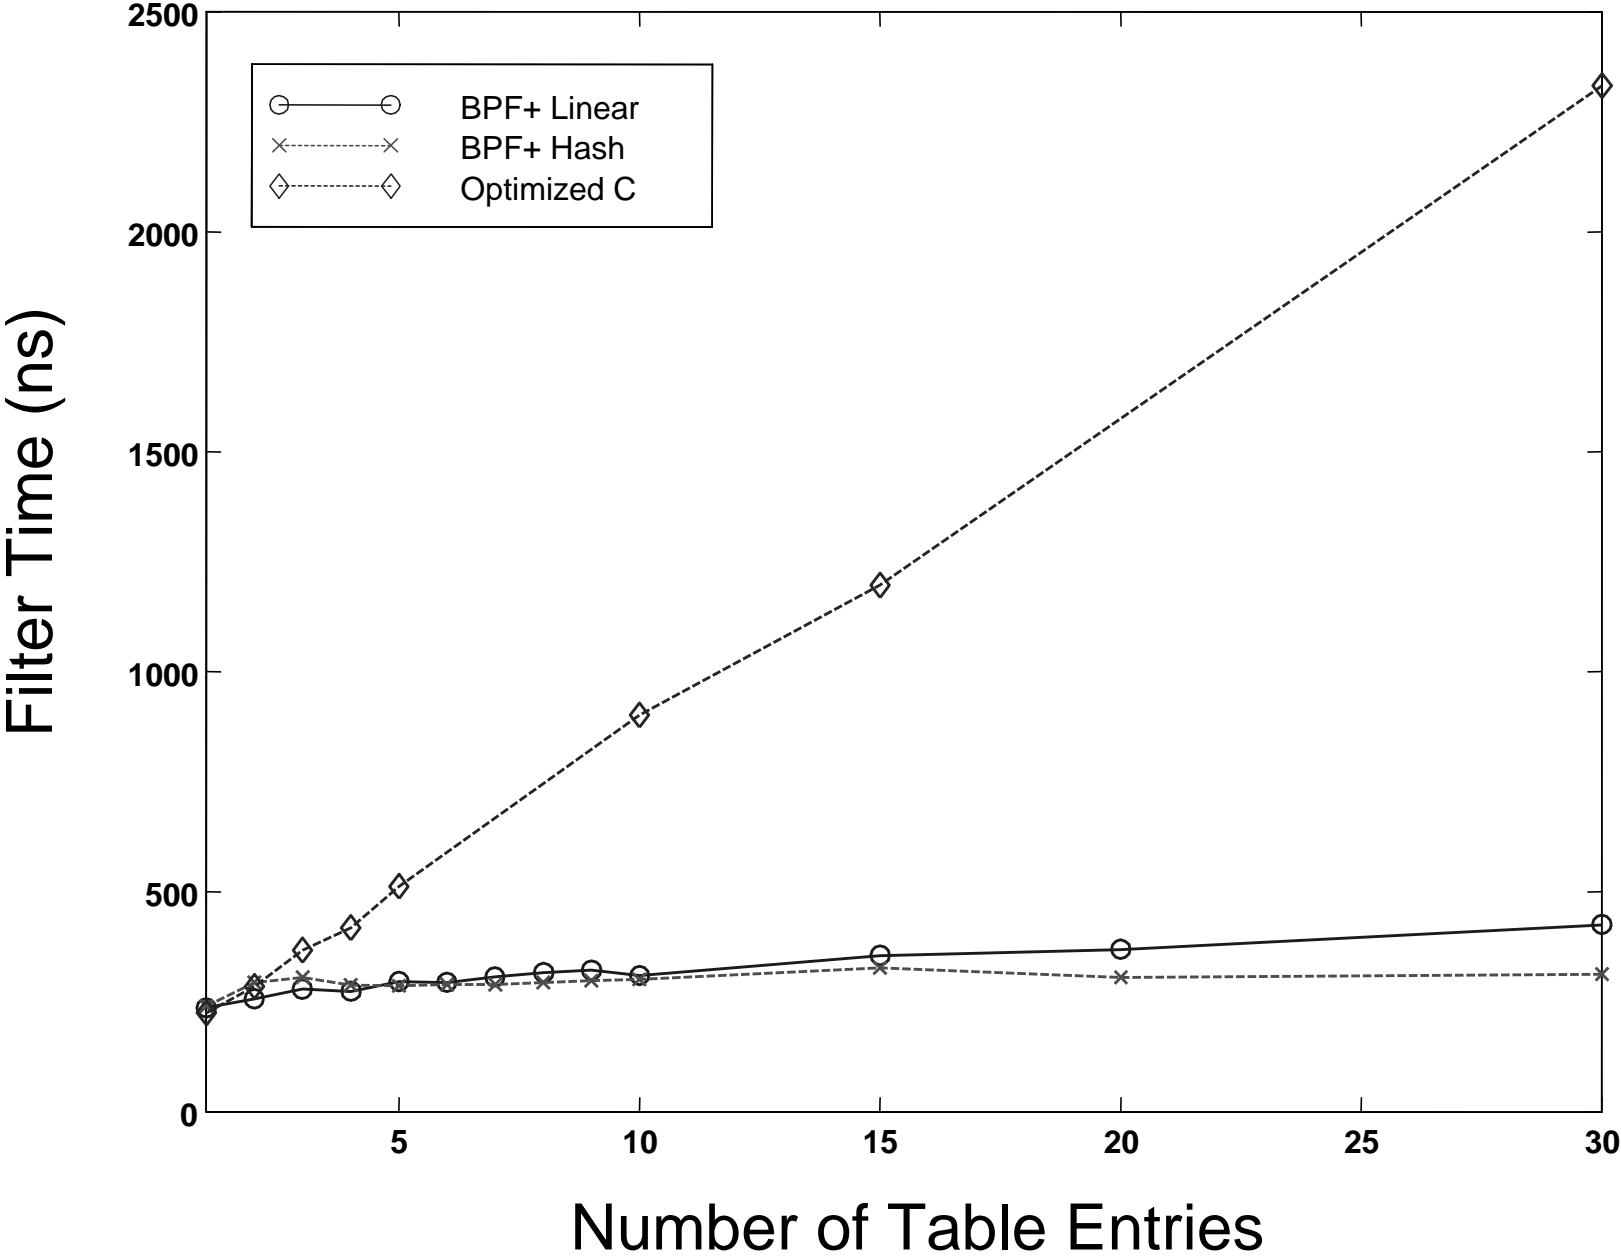
Definition: A tool for selecting packets from a packet stream using a programmable selection criterion.

Example: tcp and source net 128.32/16
and destination host web.mit.edu and port 80

Domain-Specific Optimization

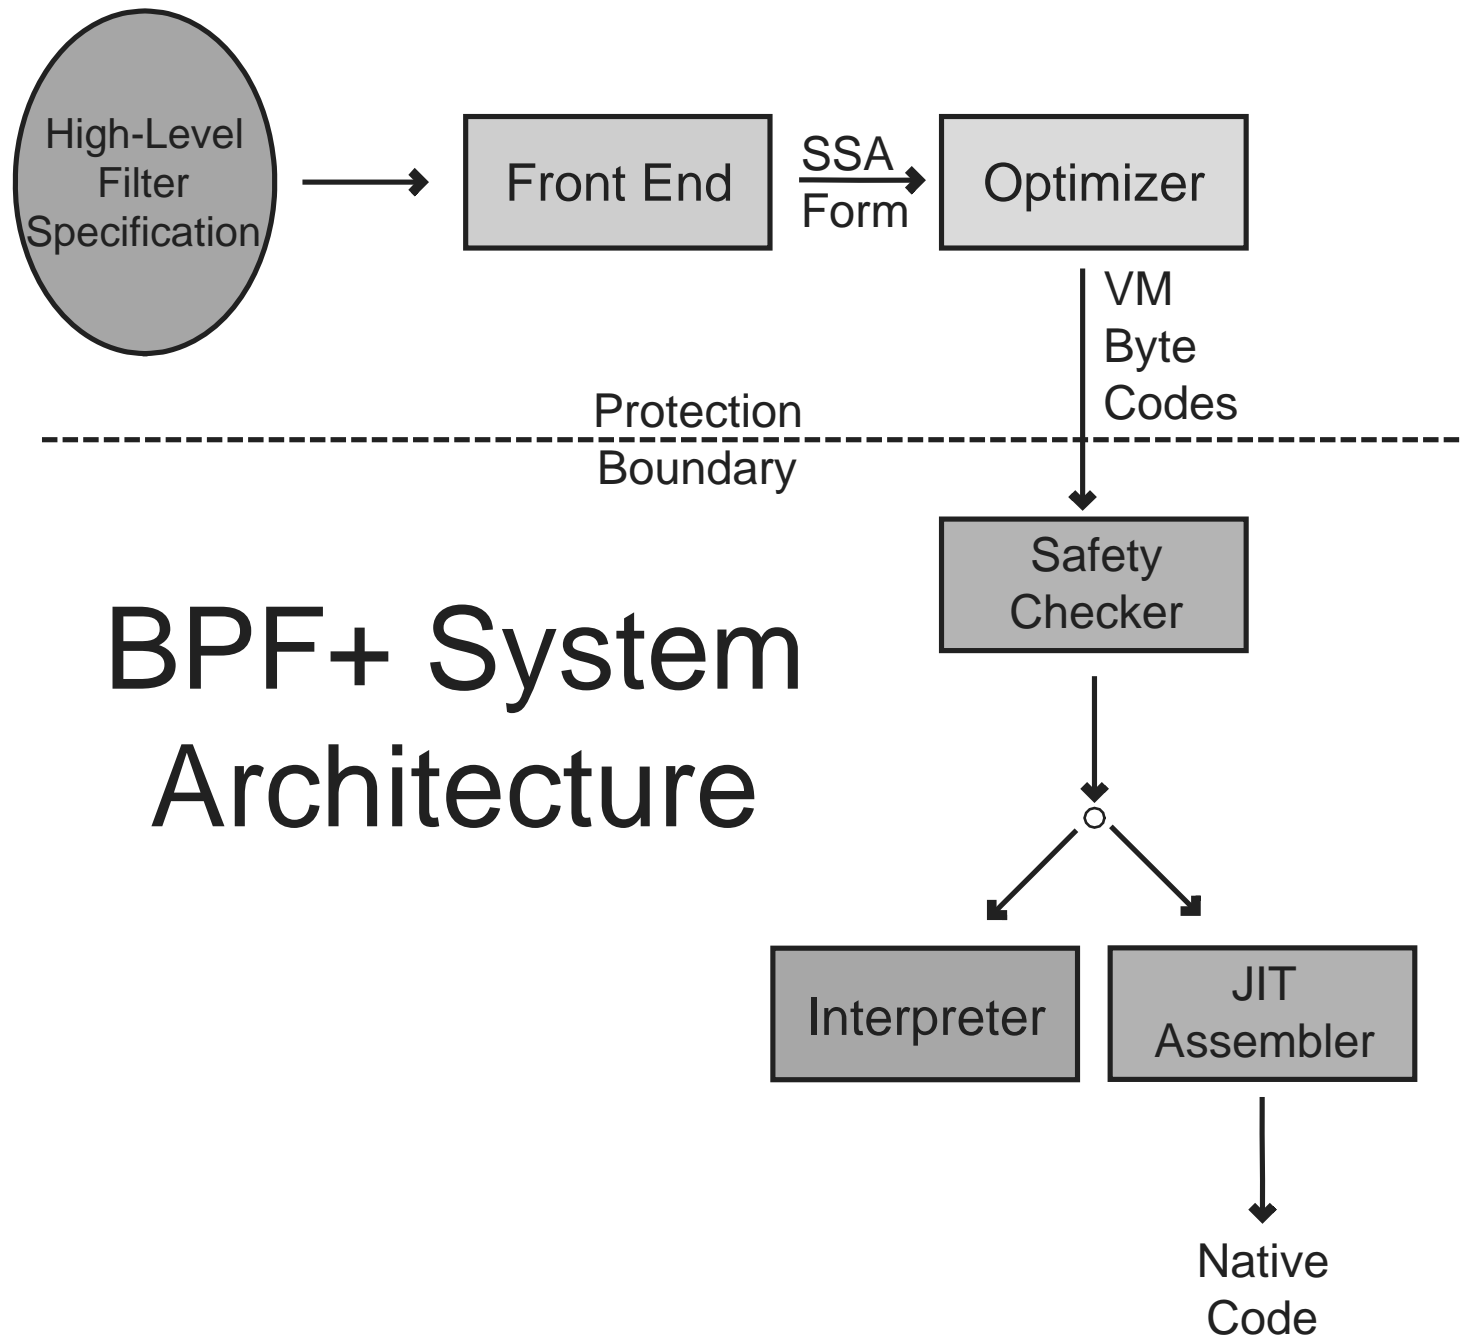
- Tune traditional compiler optimizations for the application of packet filtering.
- Affords simpler analyses and *much* more effective optimization.
 - Packet filters are DAGs
 - Engenders linear-time algorithms
- Combine with Just-In-Time (JIT) assembly.

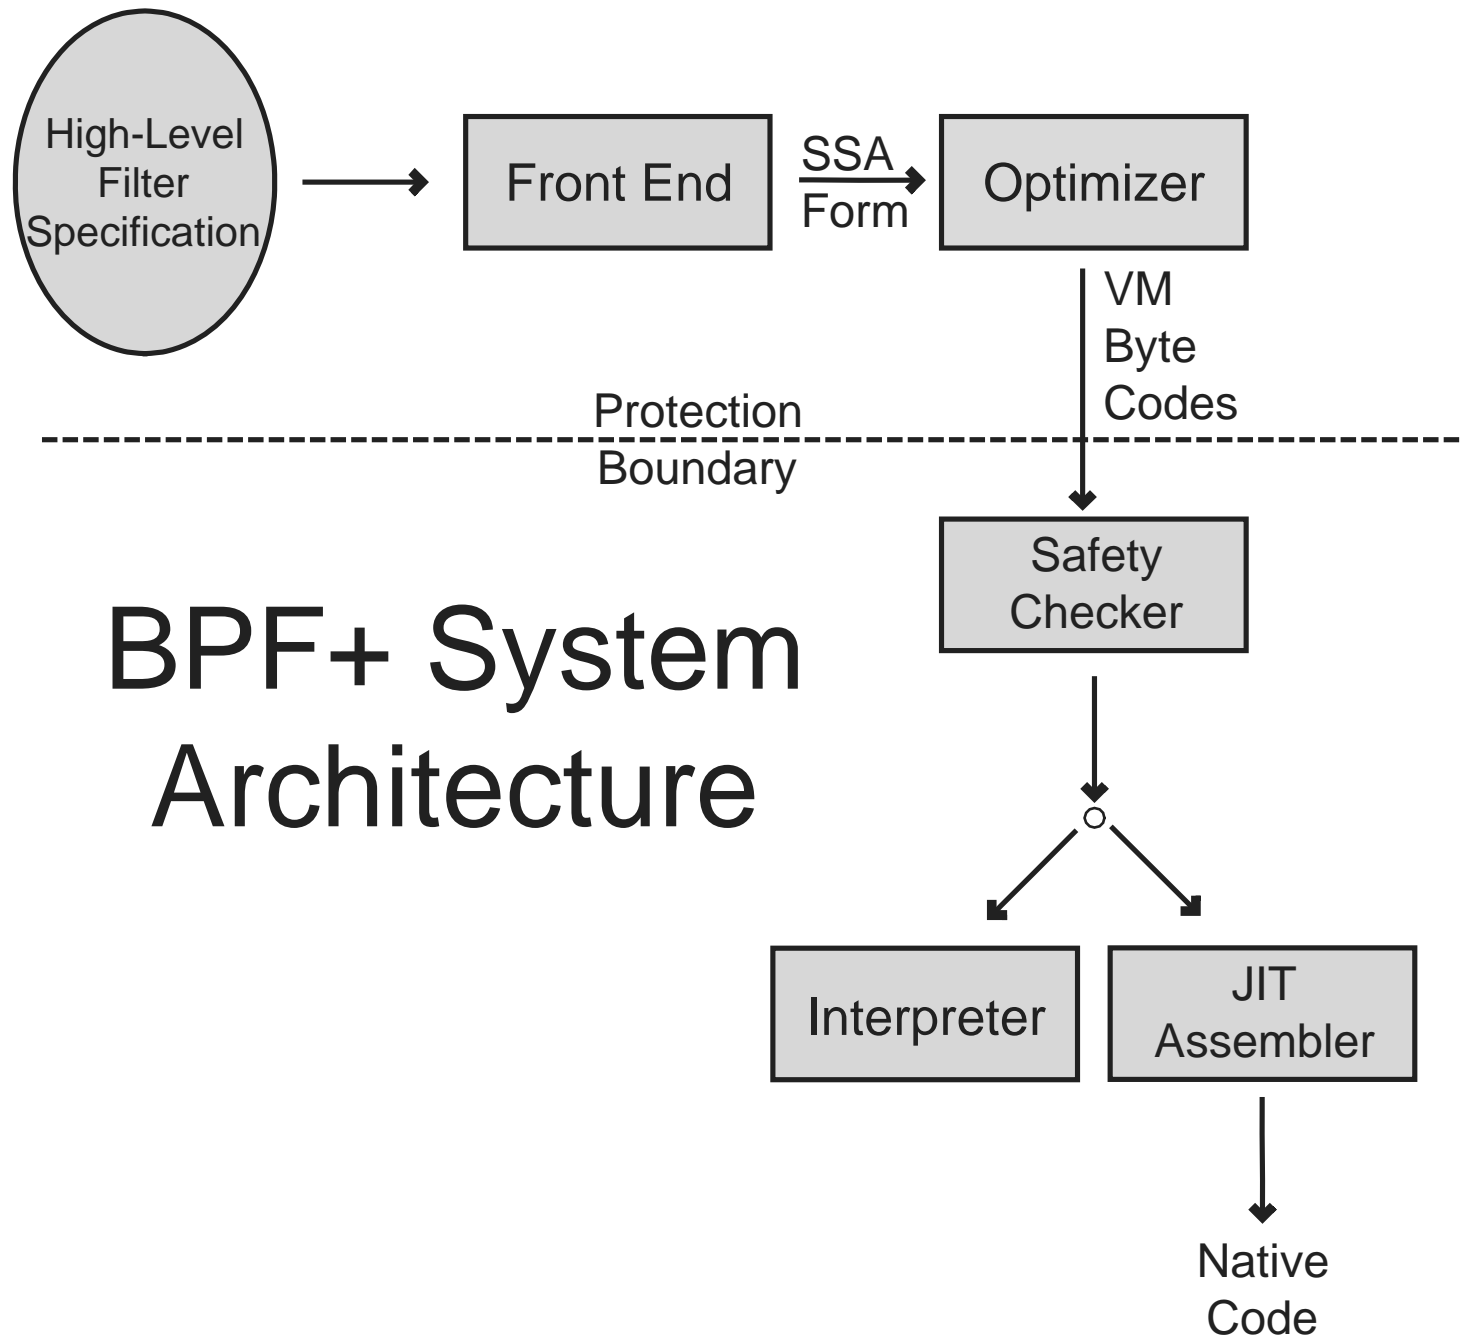
Comparison of C Optimizer to BPF+ Optimizer



Related Work

- Virtual Machine Models:
 - CMU/Stanford Packet Filter: *MRA87*
 - BPF: Berkeley Packet Filter: *MJ93*
 - tcpdump and libpcap
 - MPF: *YBMEM94*
- Exploit Filter Structure:
 - PathFinder: *BGPP94*
 - DPF: *EK96*
 - Multi-dimensional Range Matching: *LS98*
 - Grid of Tries: *SVSW98*
- High-Level Approach:
 - LR Parsing: *JC96*





Compiler Optimizations

- Leverage modern compiler technology
- Powerful intermediate form
 - Static Single Assignment (SSA)
- Three key optimizations
 1. Redundant Predicate Elimination
 2. Partial Redundancy Elimination
 3. Lookup Table Encapsulation

Redundant Predicate Elimination

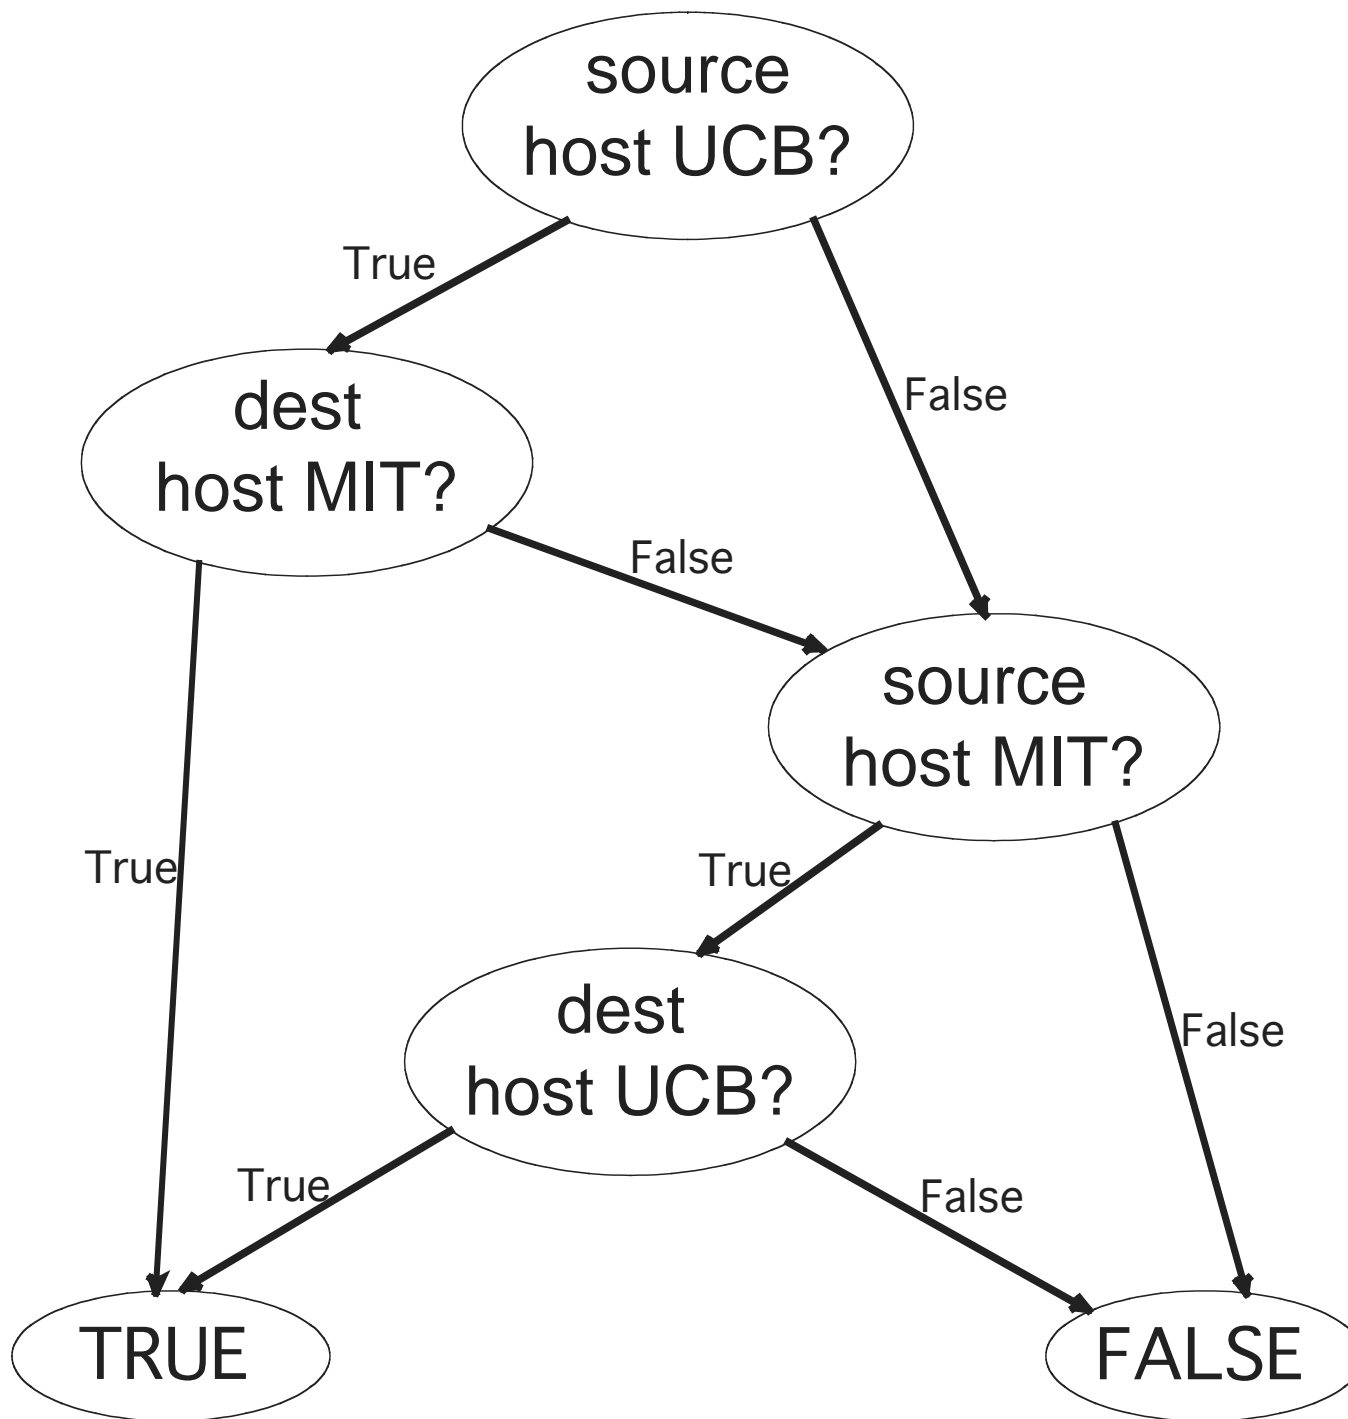
- Packet header parsing leads to many redundant predicates.
- Unlike traditional optimizations, detect redundant *edges* rather than redundant *computation*.
- Packet filter control flow graphs are edge-rich.

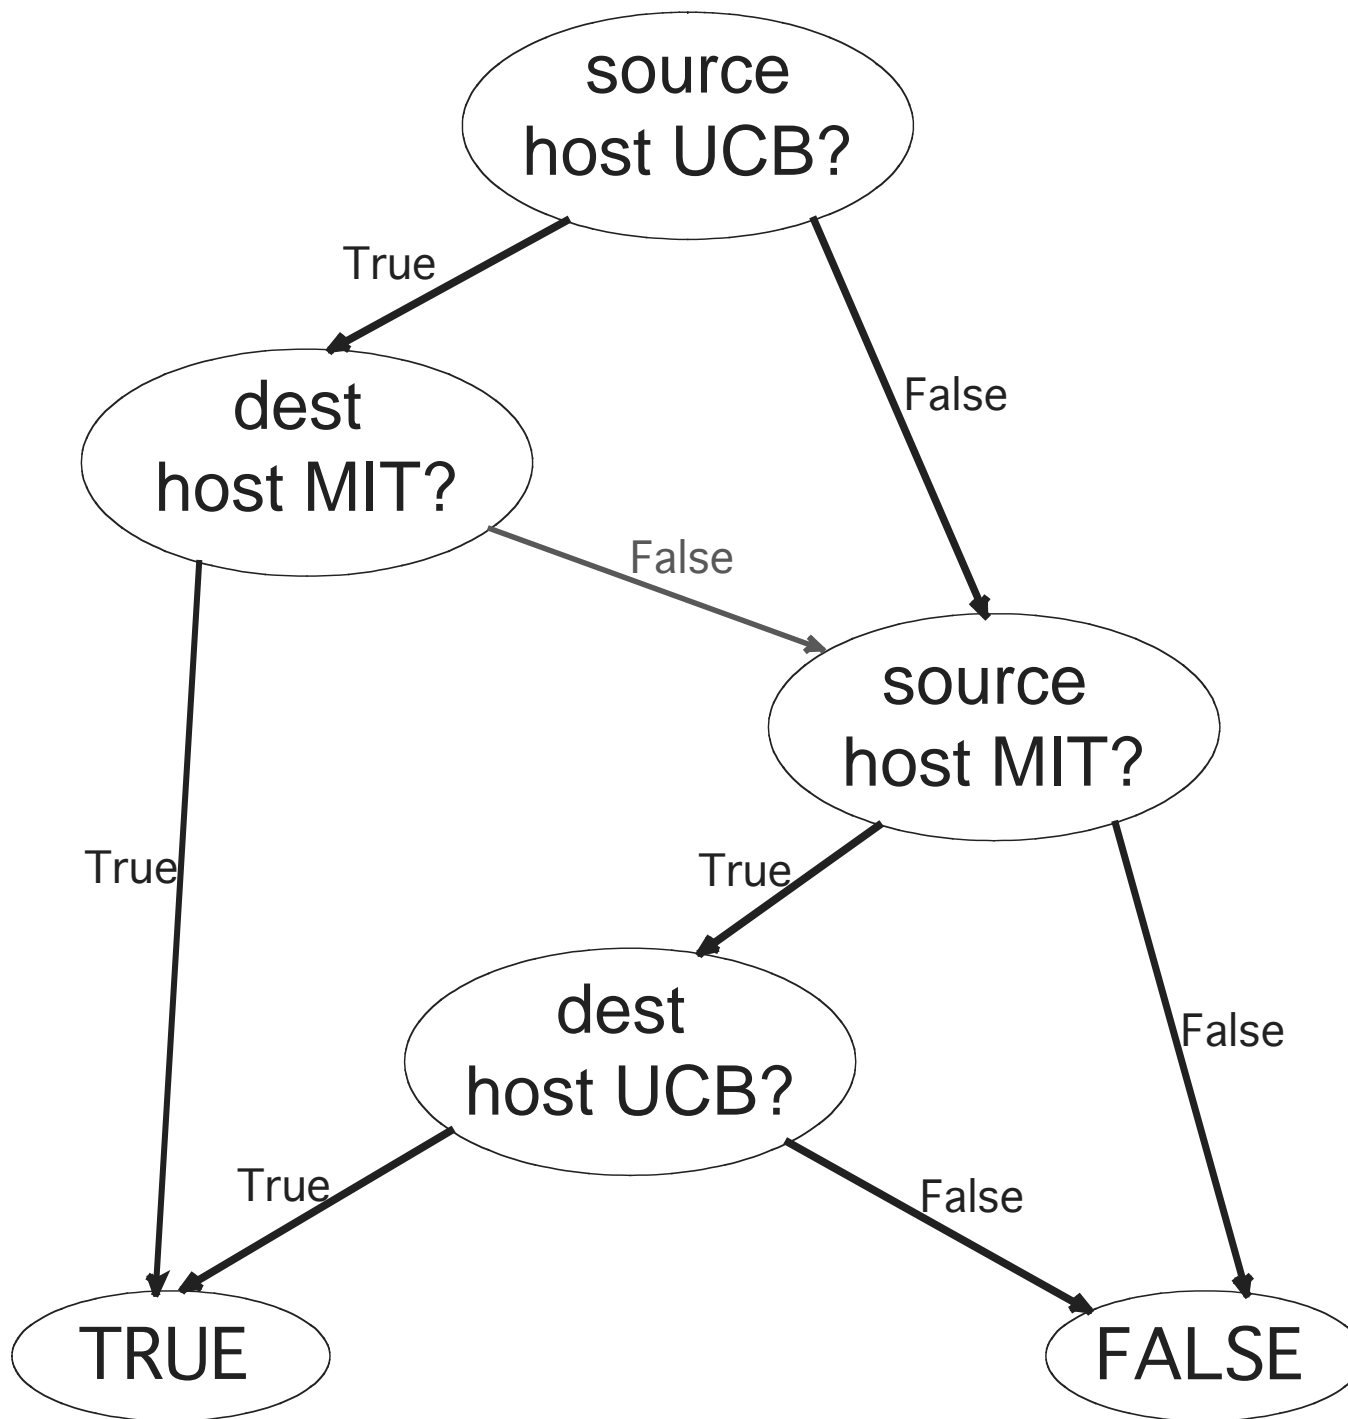
All packets between UCB and MIT

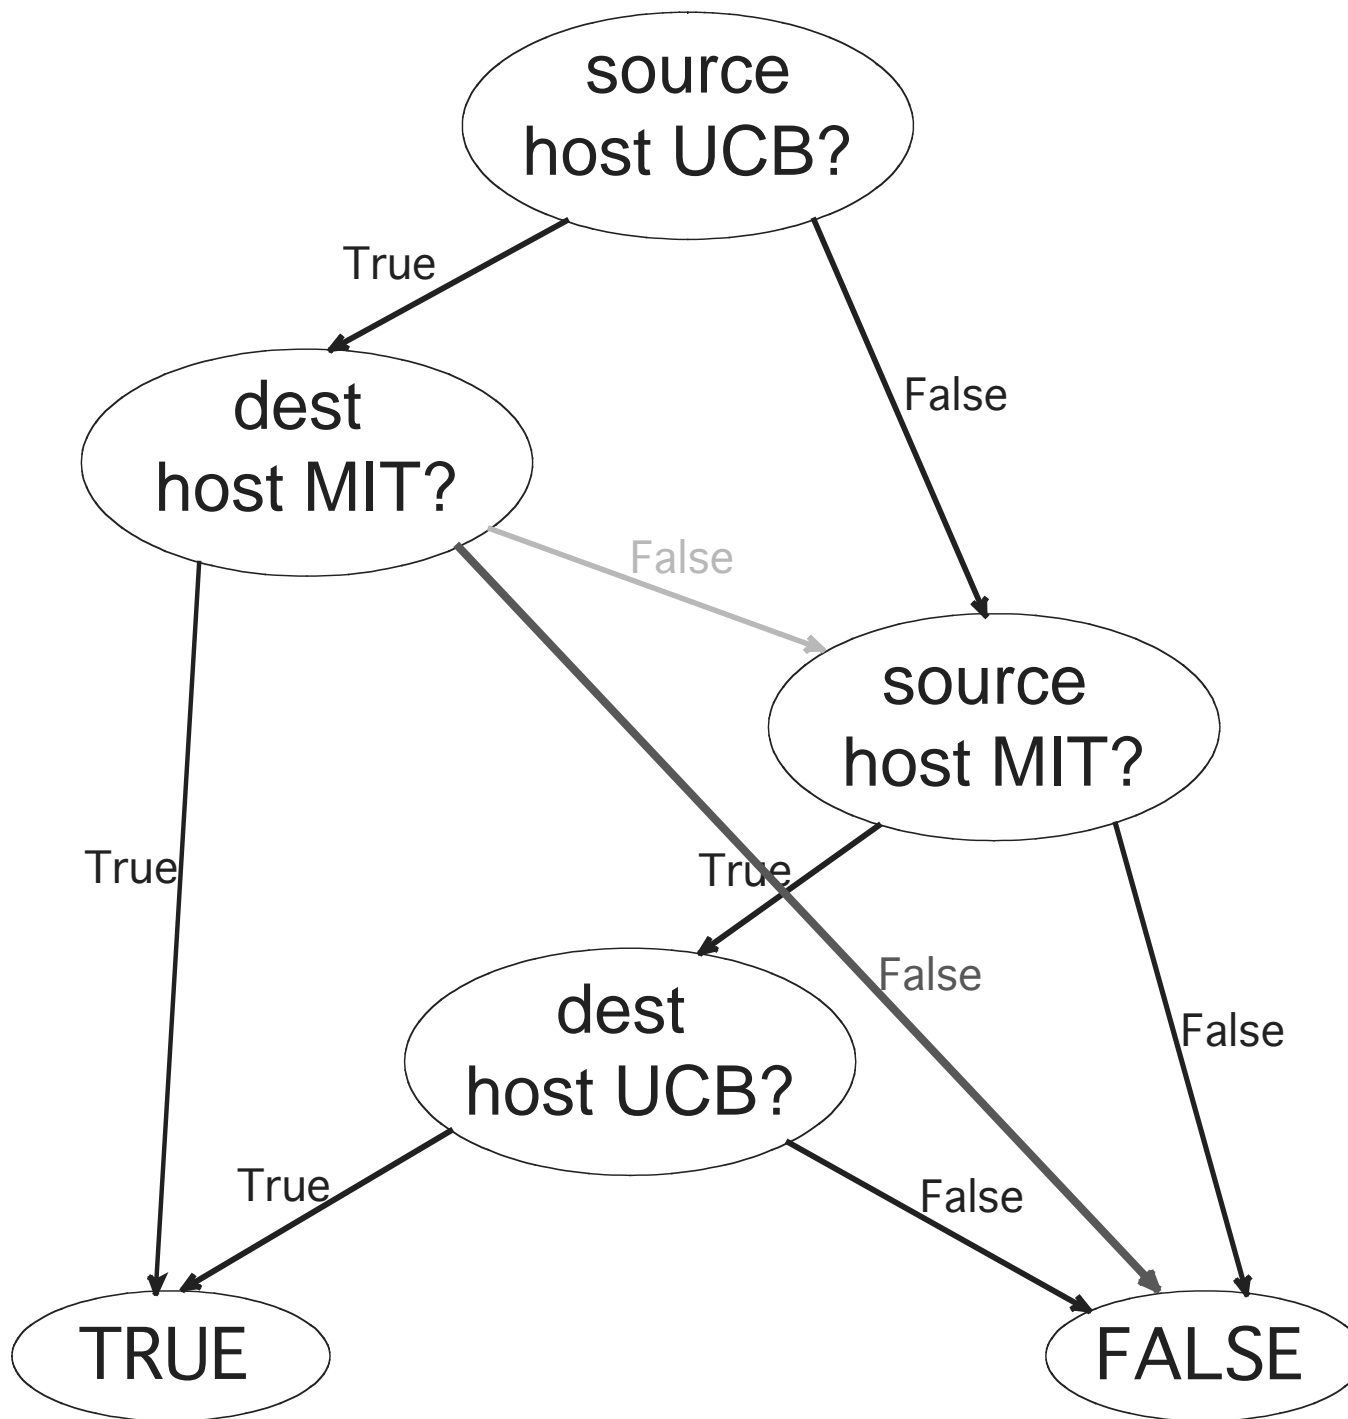
(src host UCB and dest host MIT)

or

(src host MIT and dest host UCB)





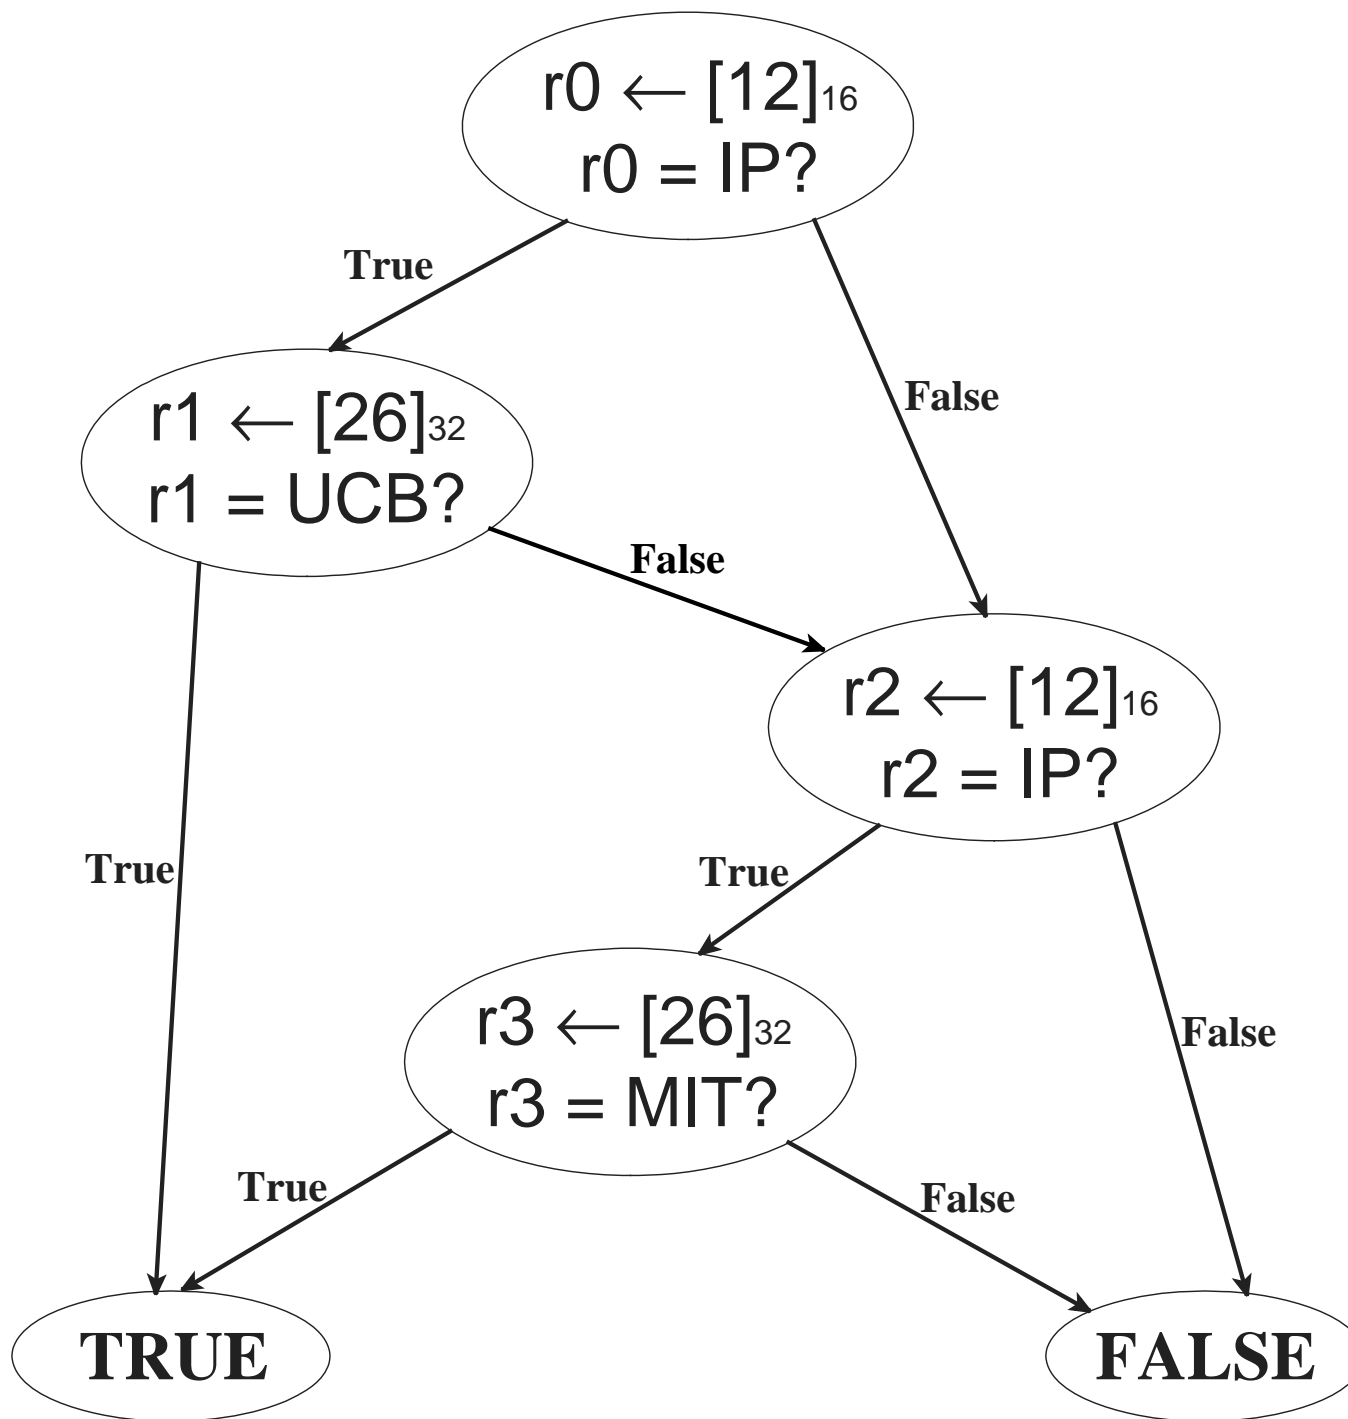


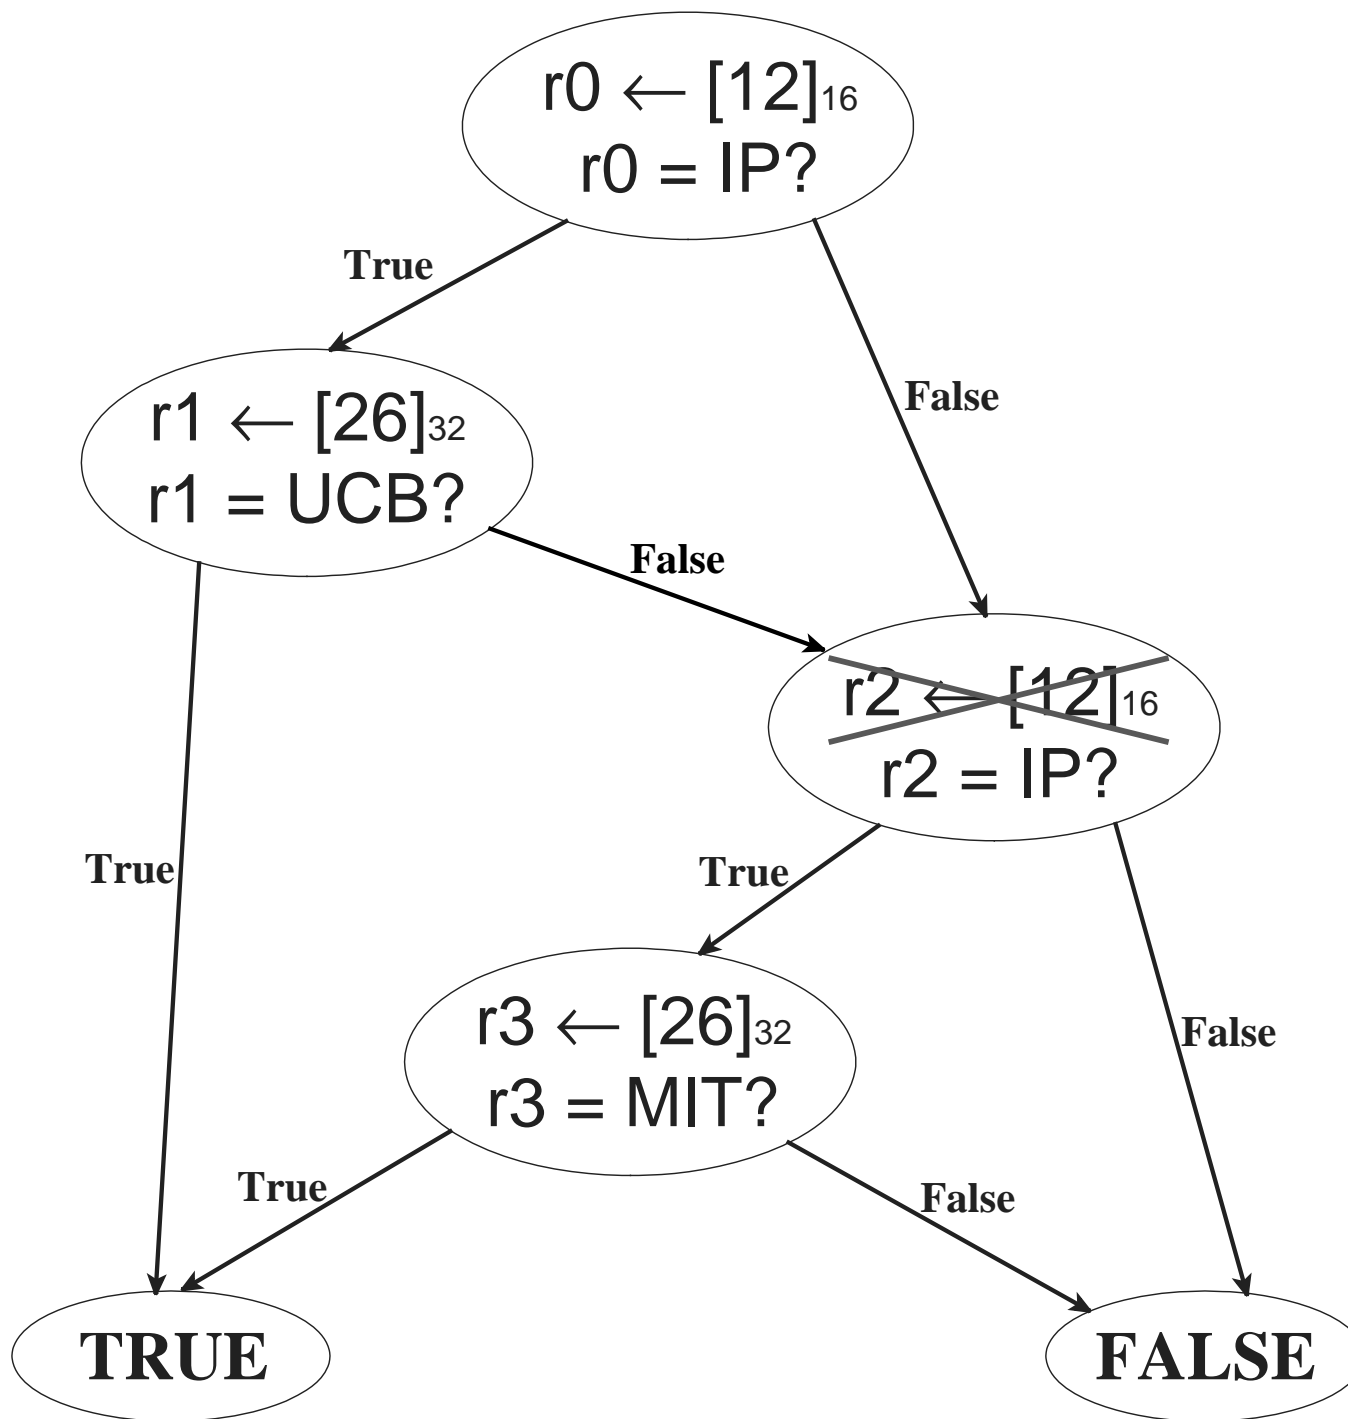
Partial Redundancy Elimination

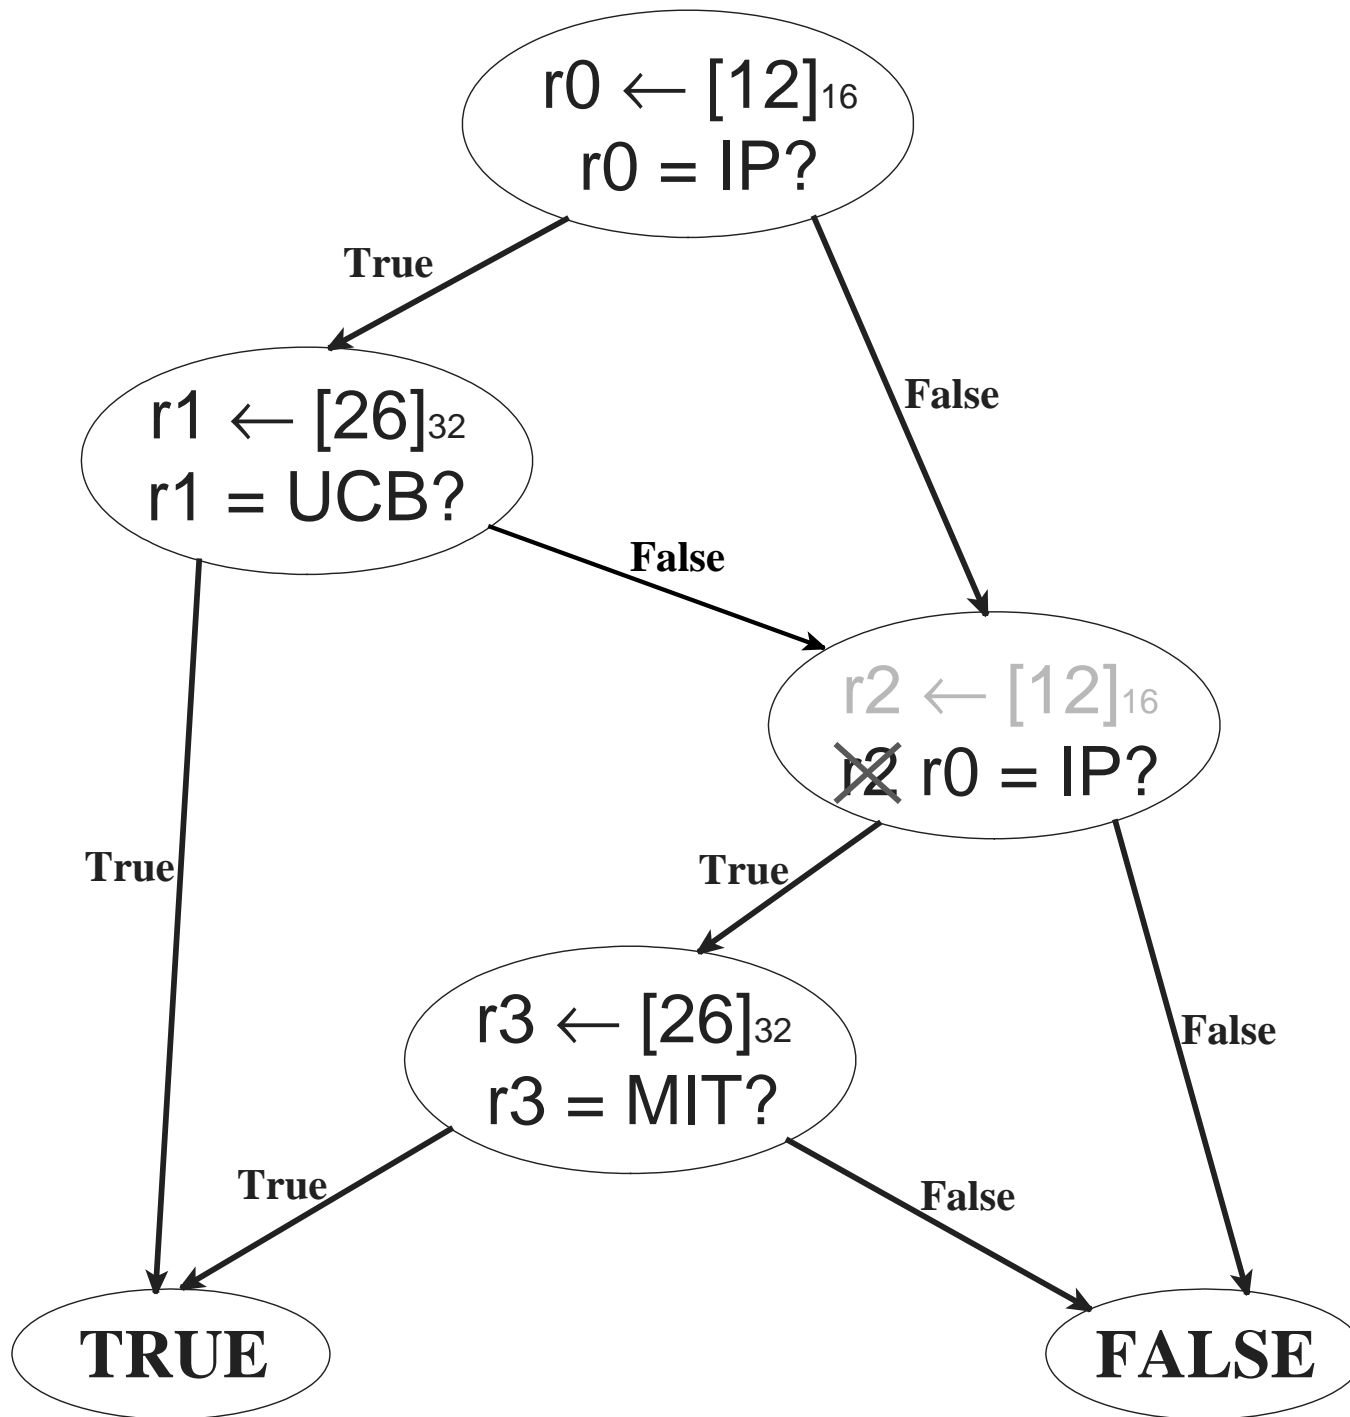
- Packet header parsing also leads to much redundant computation.

**All packets that come from
either UCB or MIT**

**src host UCB
or
src host MIT**








Putting Them Both Together

- *Combining* partial redundancy elimination with redundant predicate elimination is much more effective than using either alone.

(6 steps later...)



Lookup Table Encapsulation

- After early optimizations, many predicates reduced to simple field comparisons.
- Use analysis to discover opportunities for moving into lookup tables.
- Implementation may use linear search, binary search, hash lookup or any combination of the three.

**All packets that come from
UCB, MIT or CMU**

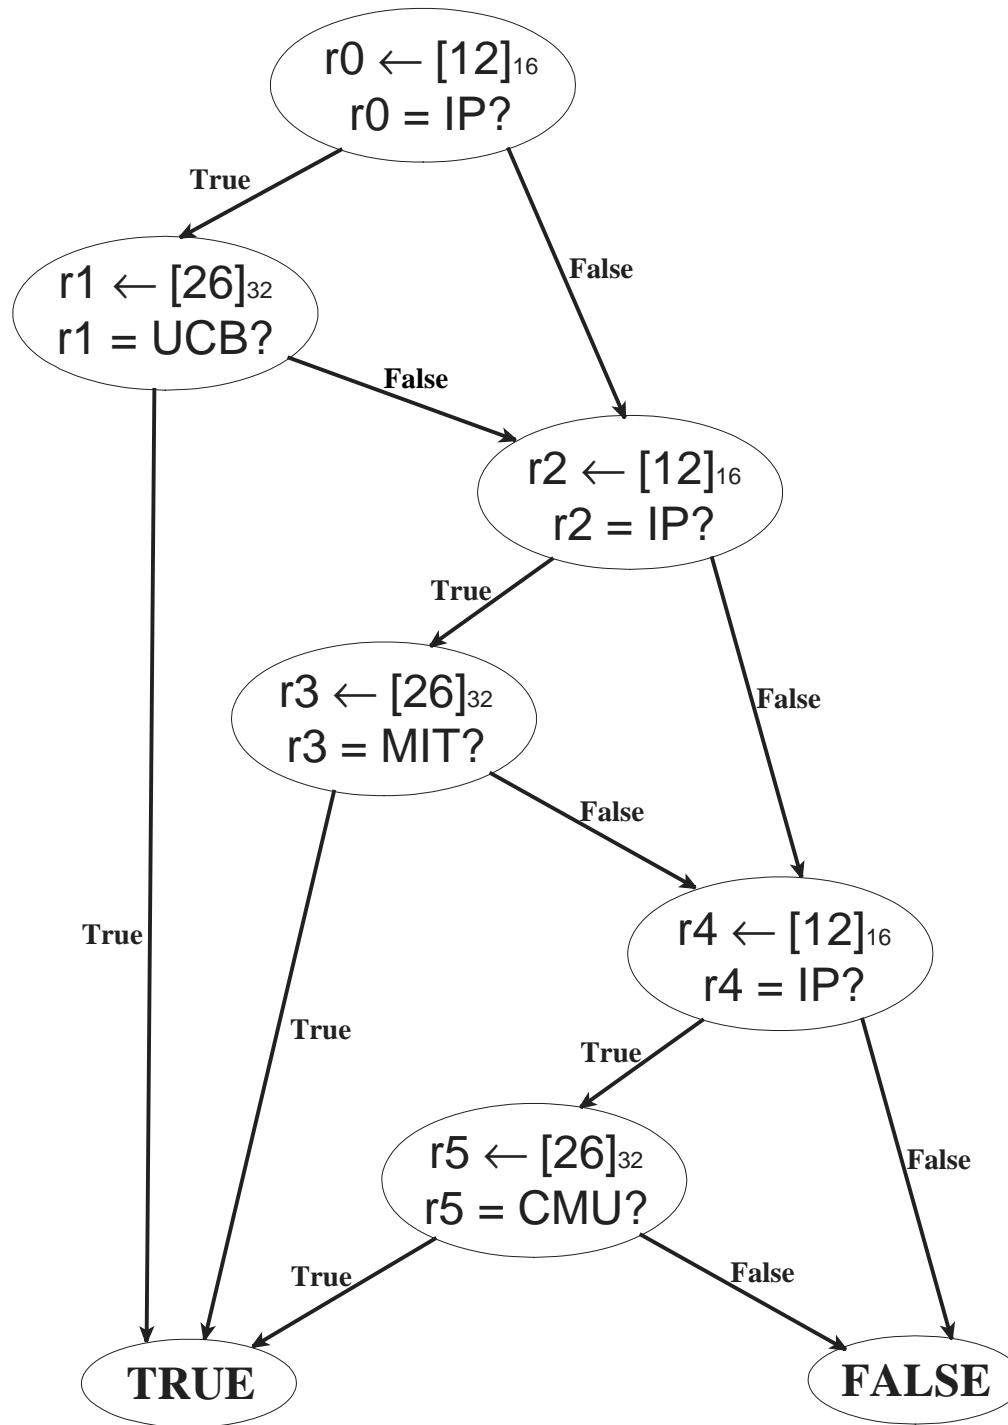
src host UCB

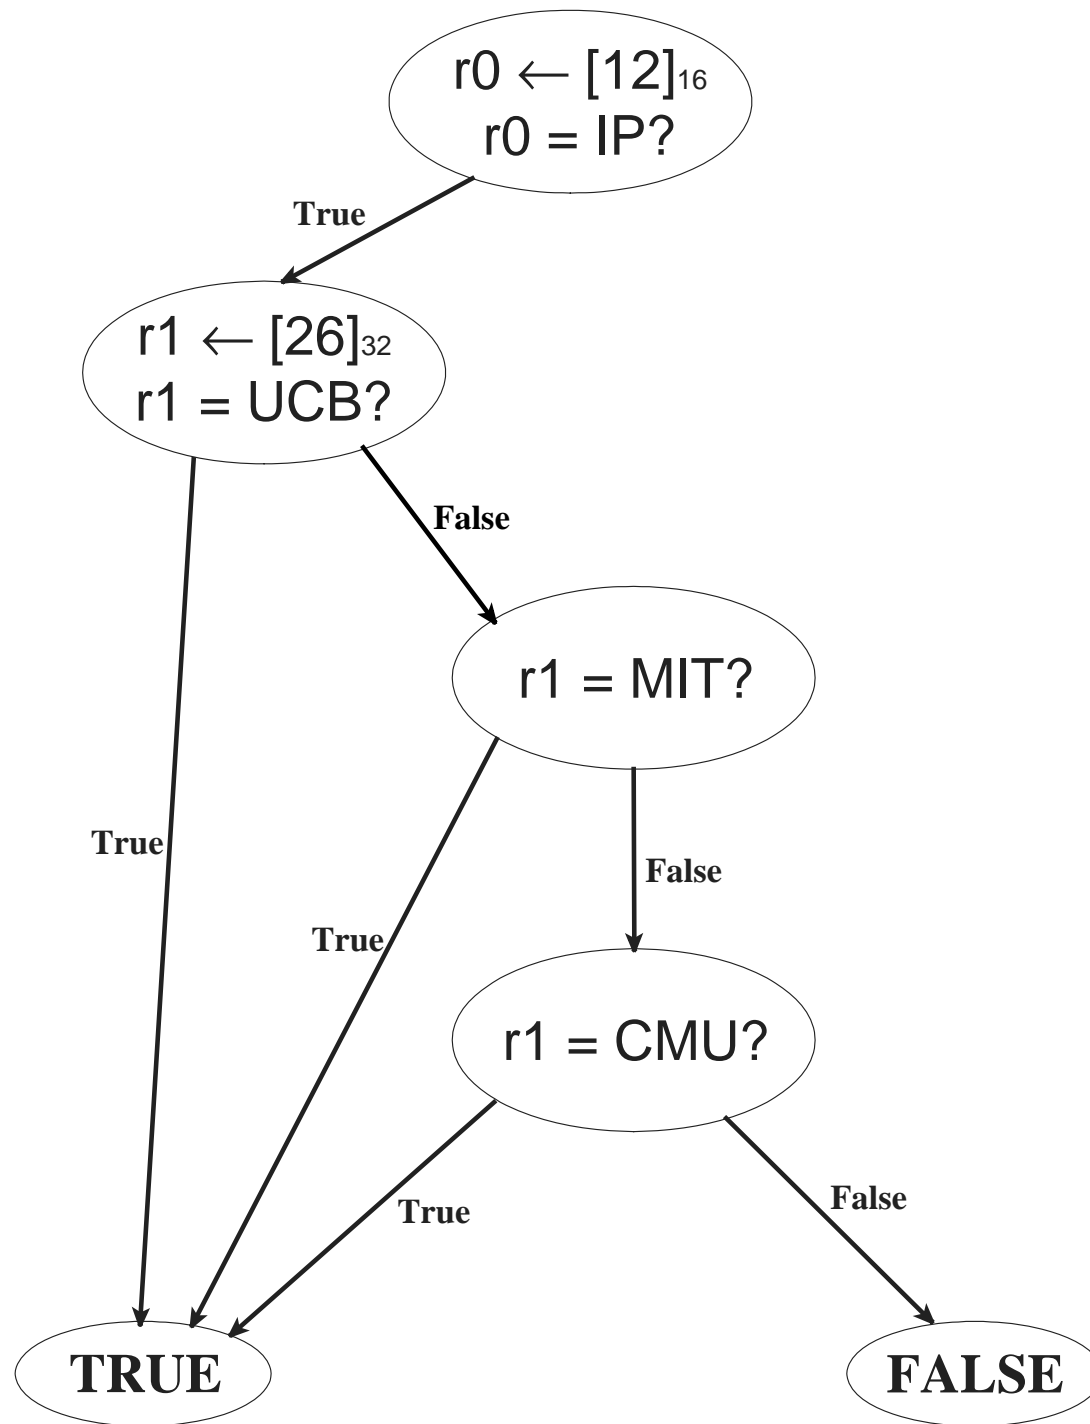
or

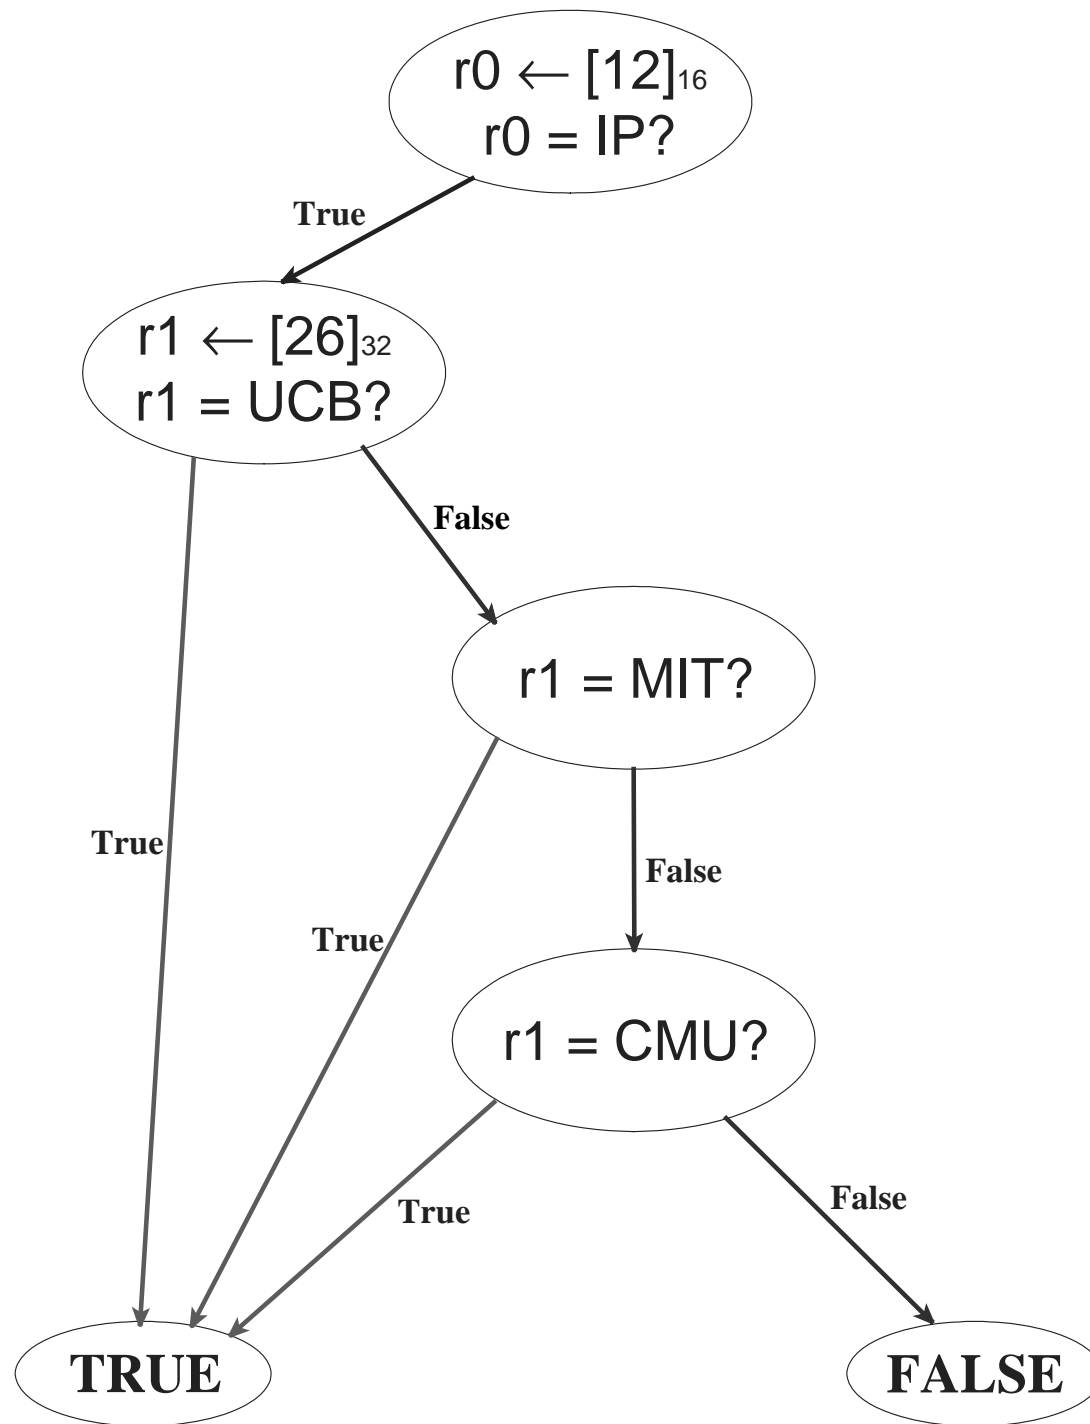
src host MIT


or

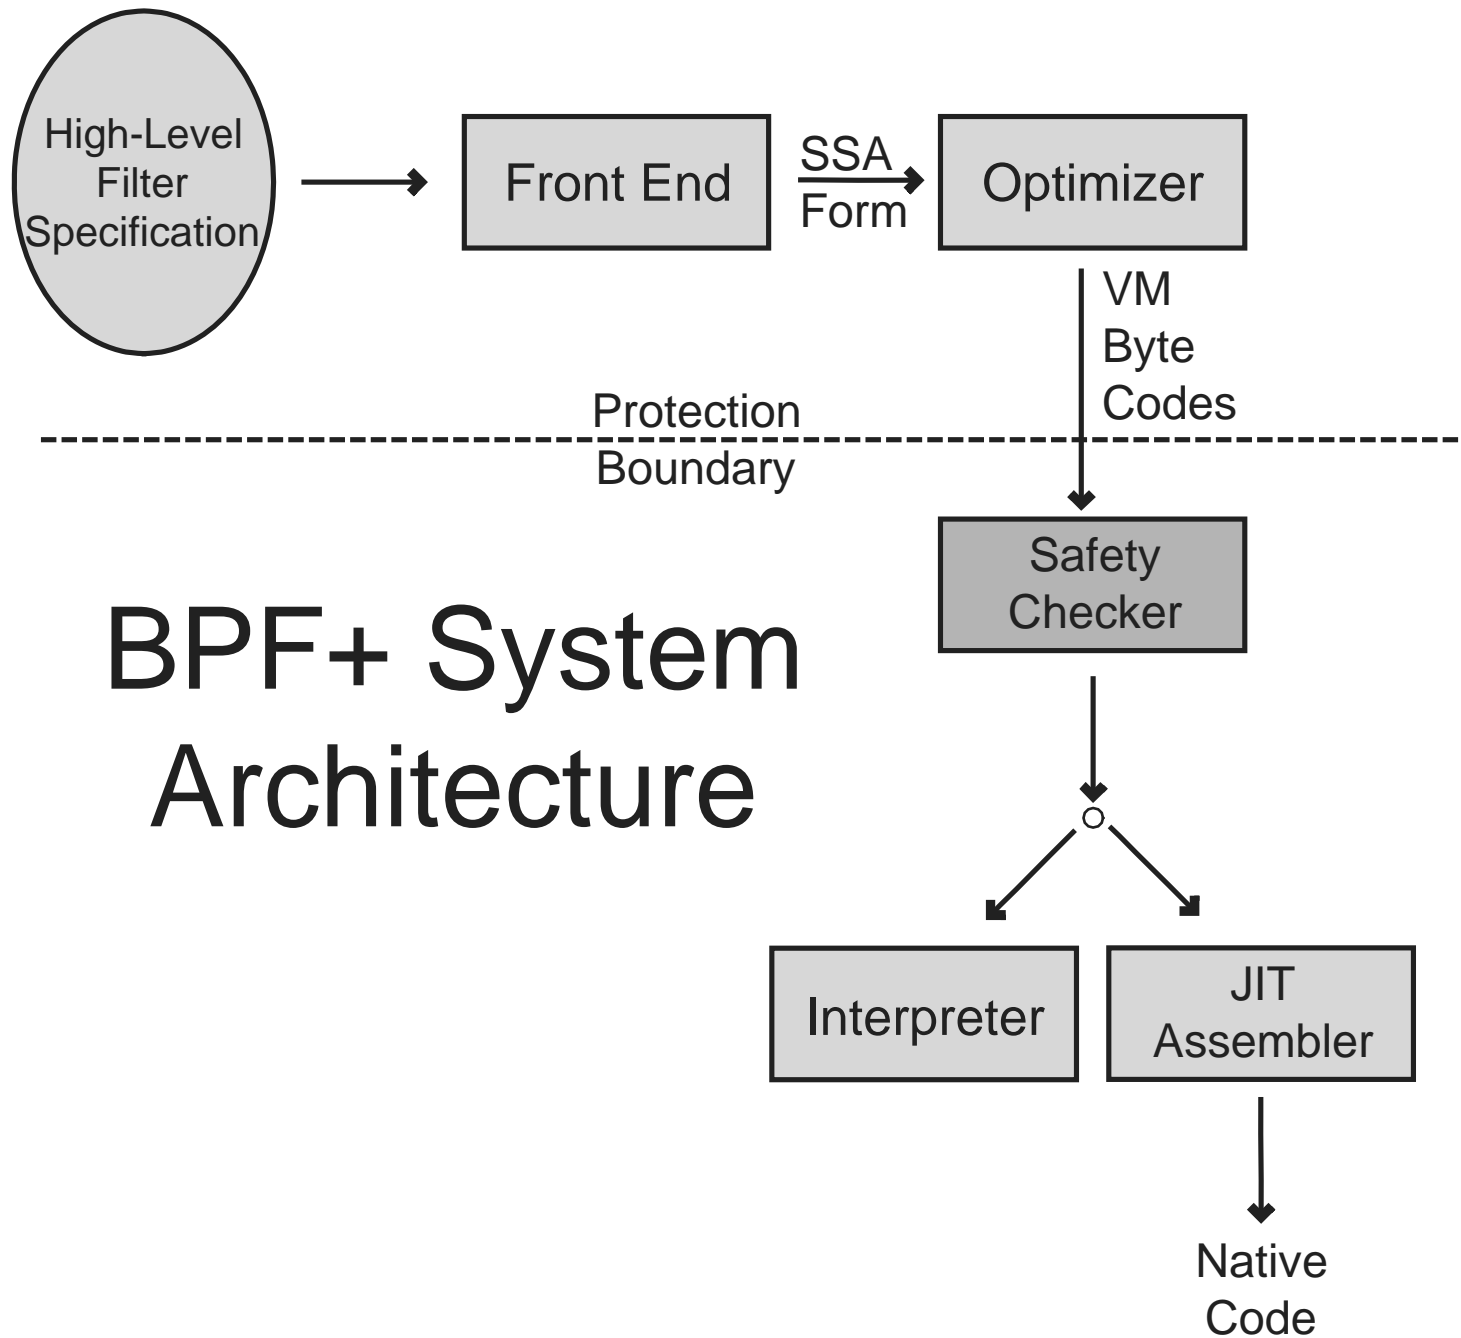
src host CMU





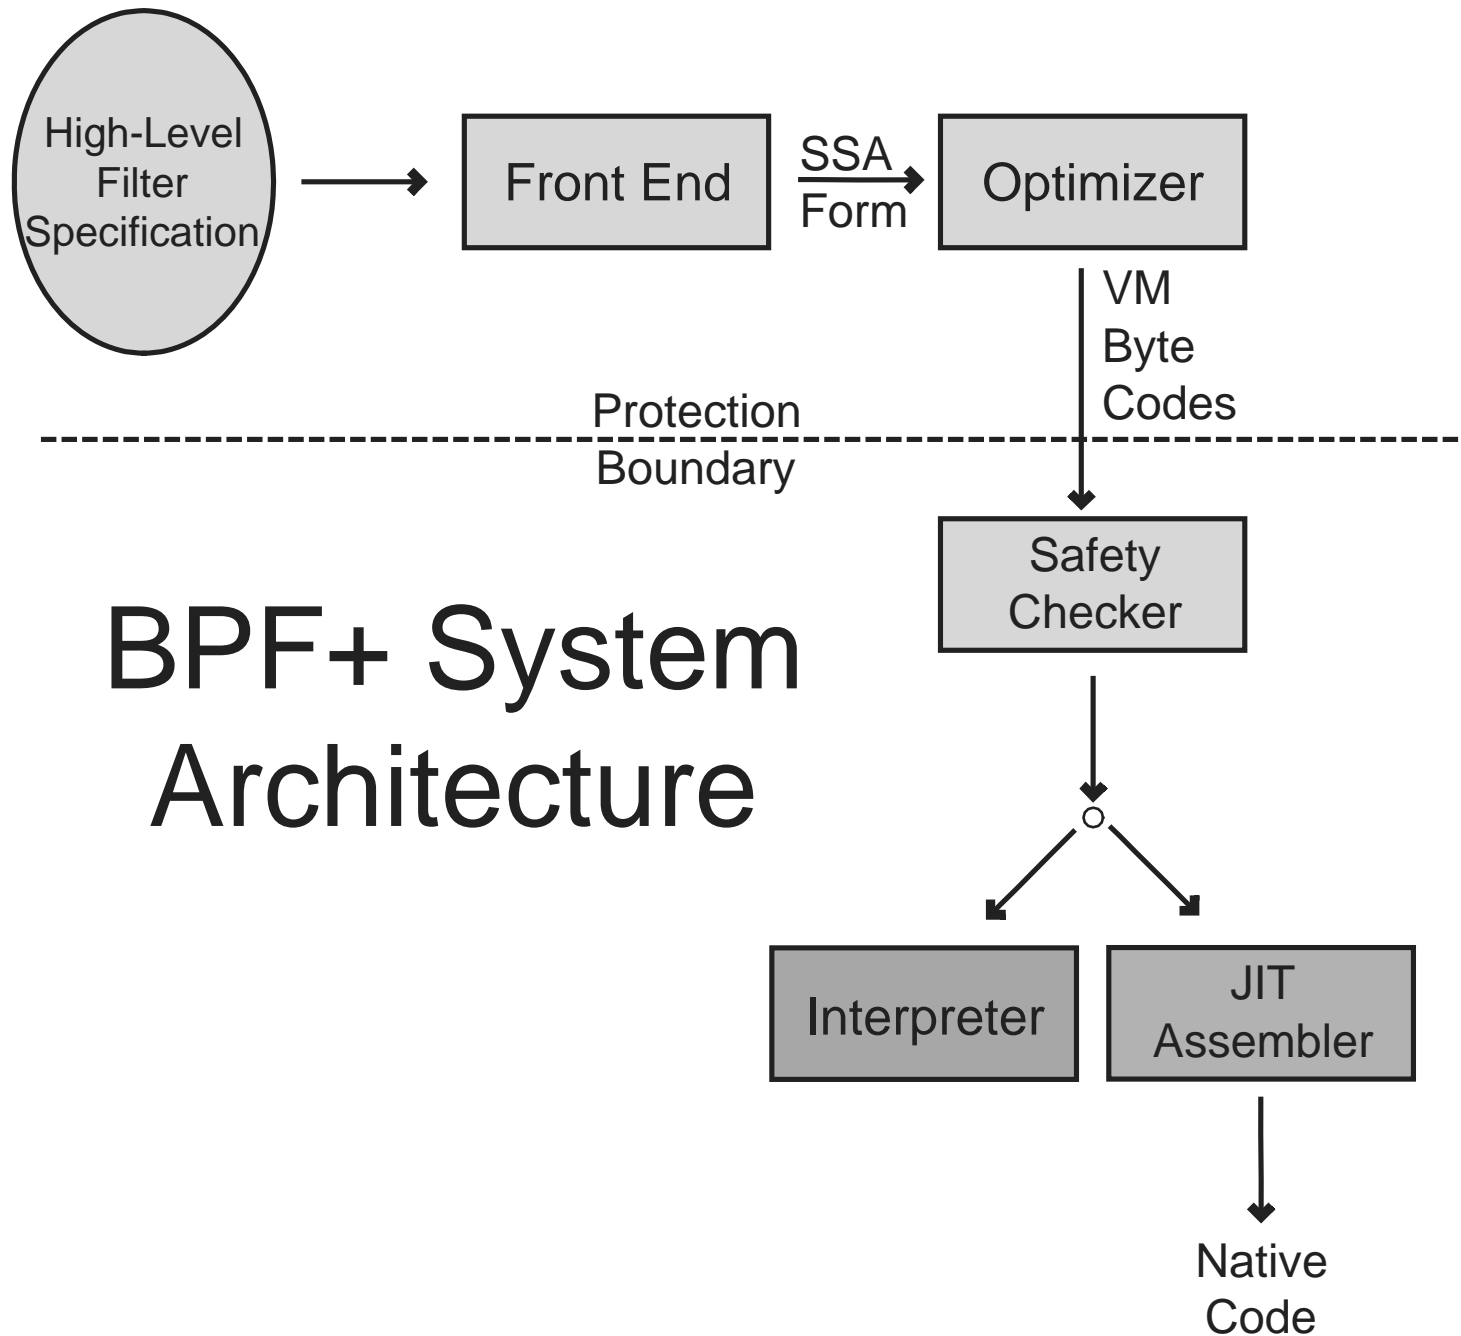






Filter Safety Must Be Verified

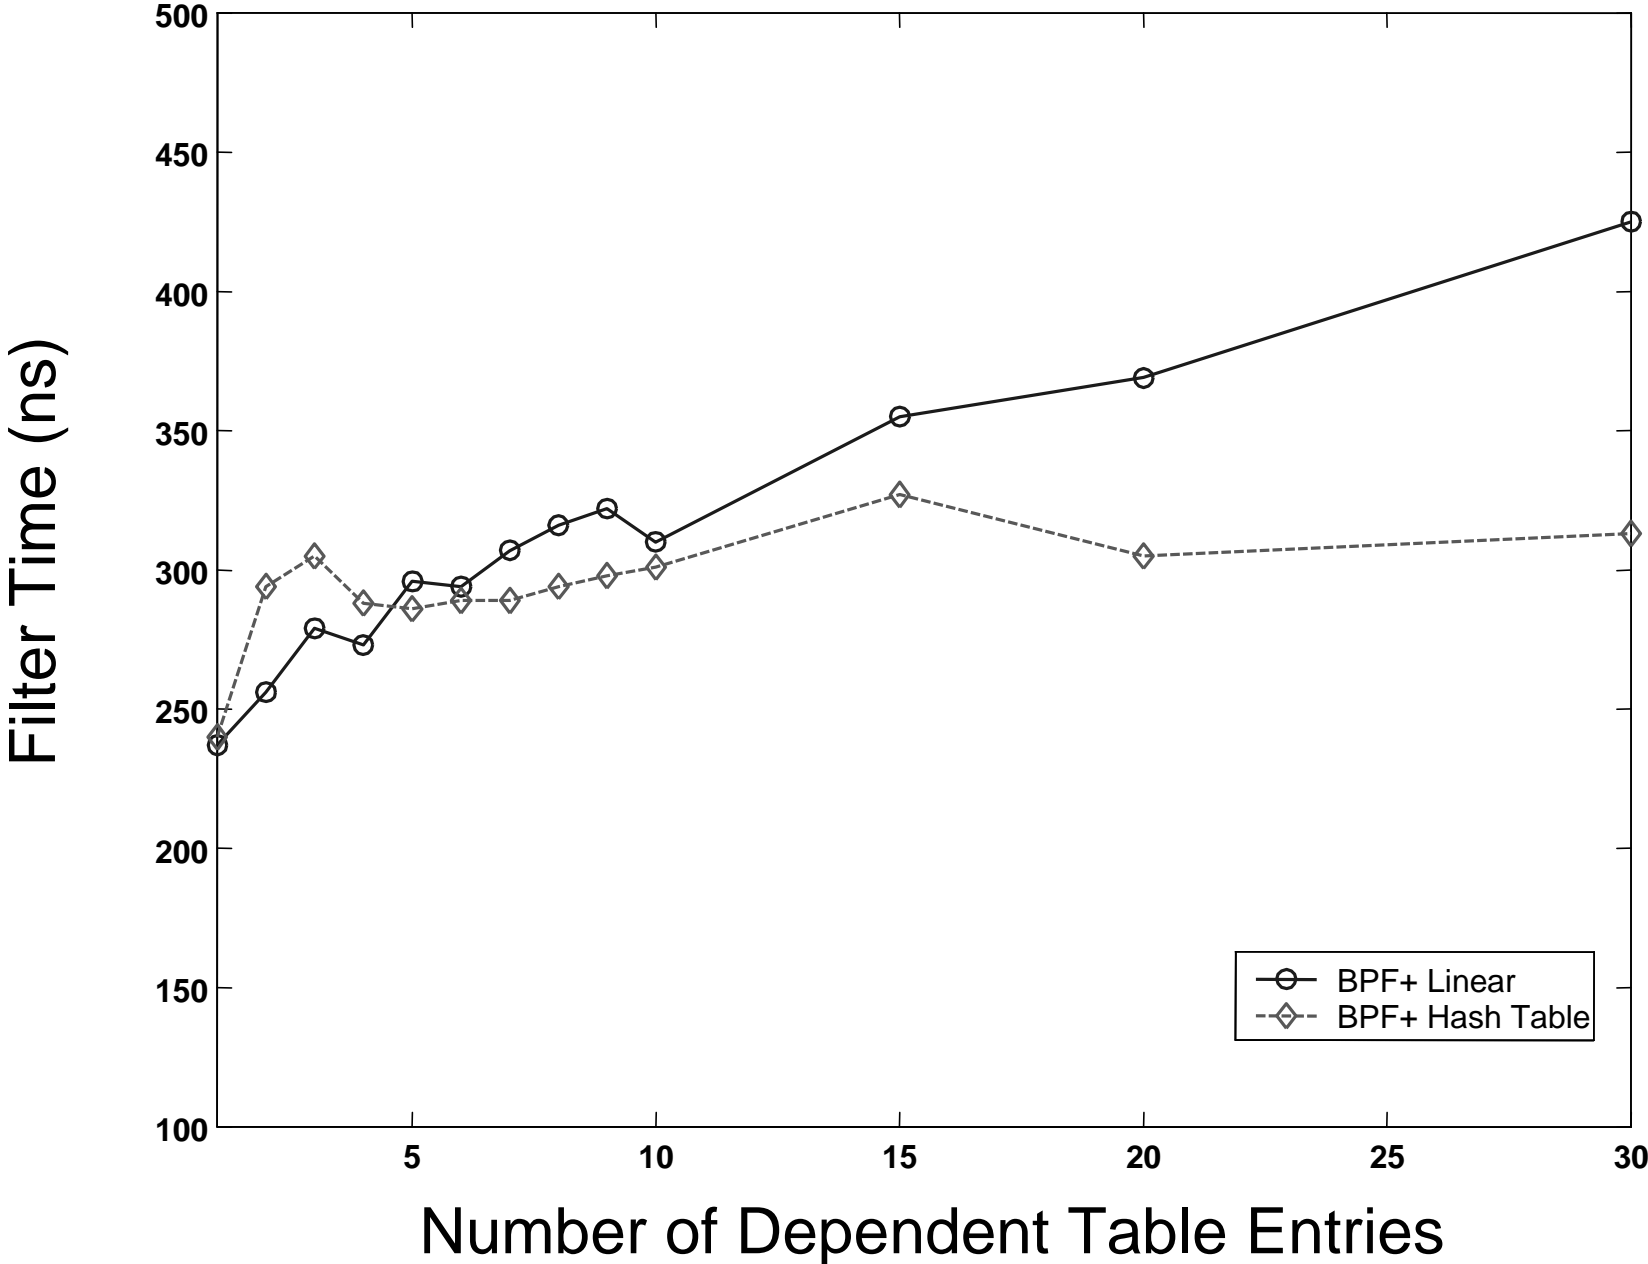
- All bytecodes are valid.
- Jump targets are valid.
- No loops.
- All paths terminate with a **return** instruction.
- No out-of-bounds reads or writes.



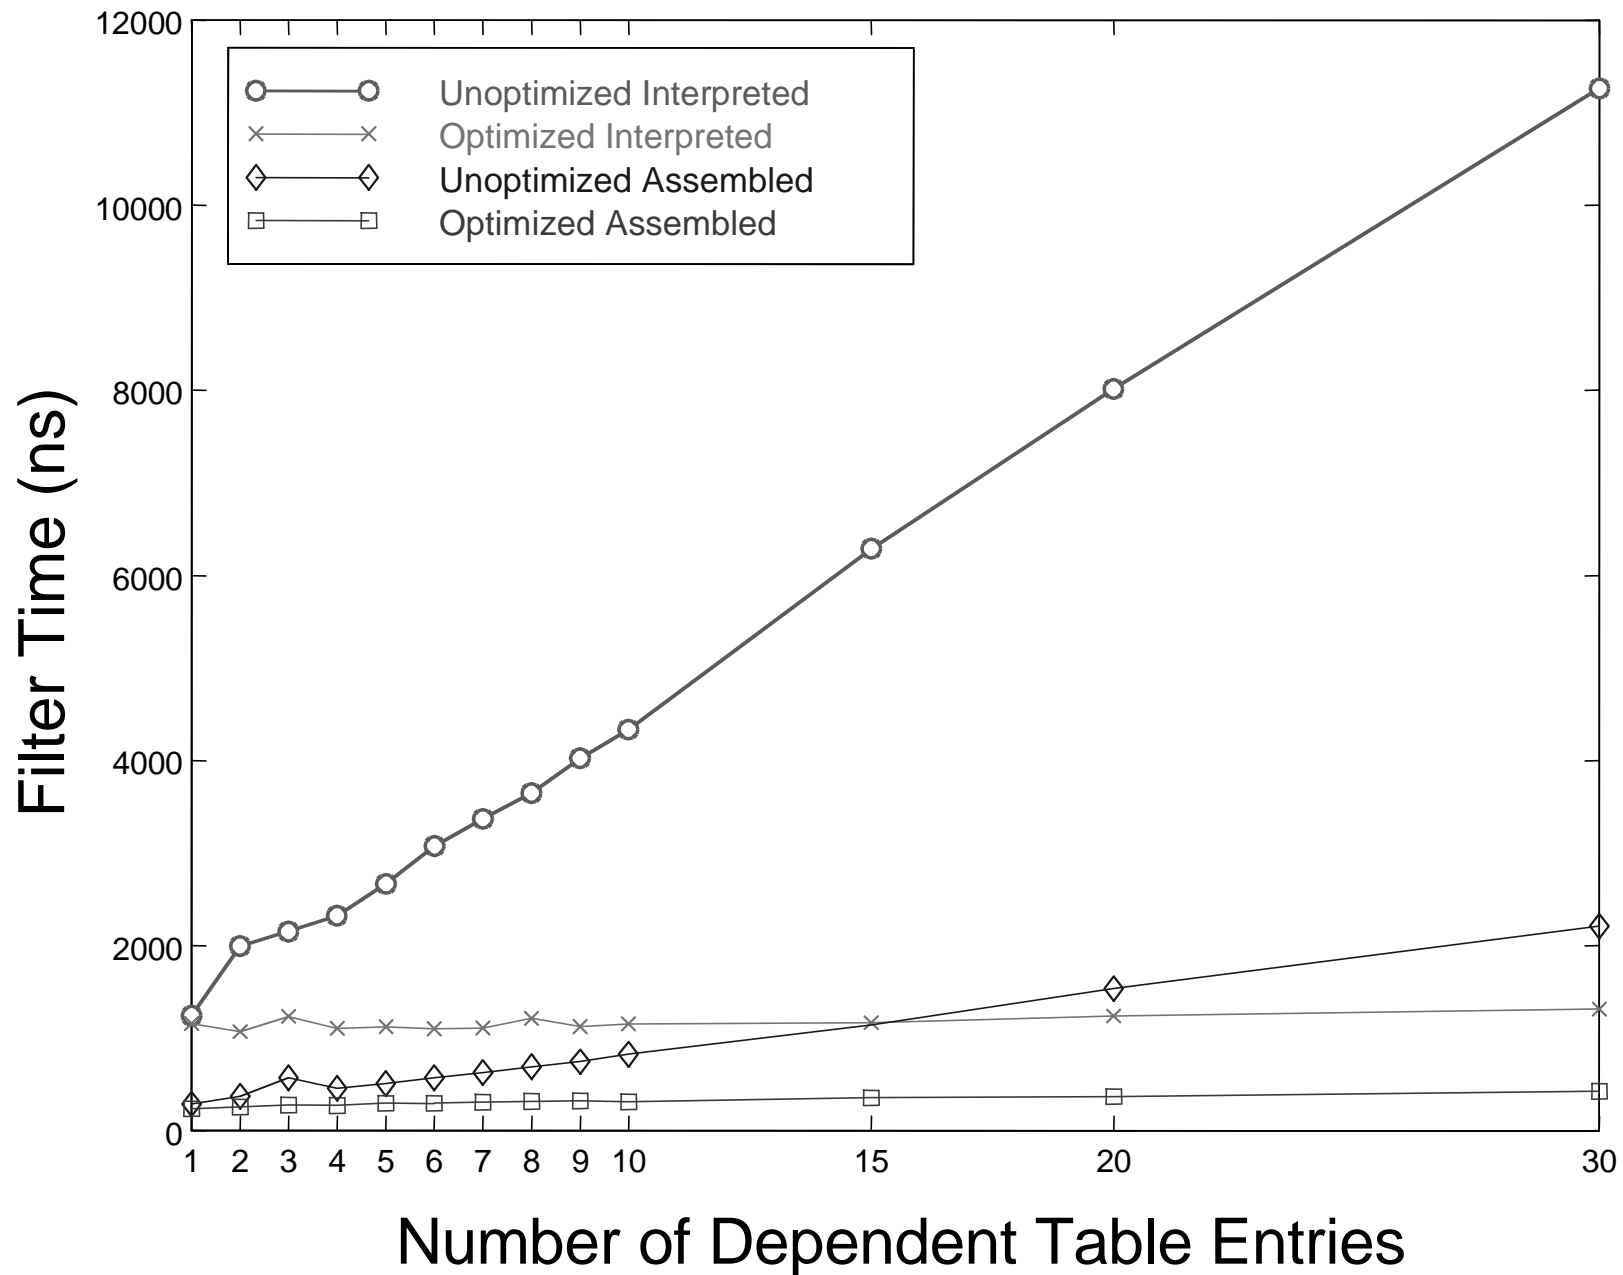
Performance Tests

- Two types of predicate expressions
 - Dependent (src host (UCB or MIT or CMU)
 - Independent (i.e. TCP, Port 80, dest host UCB)
- Run on Ultra 10 300 MHz UltraSPARC IIi
- Four ways to run a filter
 - Unoptimized, Interpreted
 - Optimized, Interpreted
 - Unoptimized, JIT Assembled
 - Optimized, JIT Assembled


Comparison of Linear Search to Hash Lookup




Effects of Optimization and JIT Assembly



Effects of Optimization and JIT Assembly (log scale)



Effects of Optimization and JIT Assembly on Independent Predicates (log scale)



Future Work

- More efficient table lookup representations (LS98, SVSW98)
- Better support for packet classification
- Loops
 - Proof-Carrying Code, Necula 96
- Intrusion detection
 - Online Updates

Conclusions

- Packet filters can be specified at a high-level *and* be efficiently executed.
- Key idea: Tune familiar global data-flow compiler analyses and optimizations for packet filtering.